

Machine Learning with Python

1 Part1

Most of the time, it's also possible to **convert a supervised dataset to unsupervised**

Firstly, first and foremost, let's start with importing required libraries.

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns

from sklearn import datasets
from sklearn import manifold
```

1.1 Cross-Validation

- k-fold cross-validation
- stratified k-fold cross-validation
- hold-out based validation
- leave-one-out cross-validation
- group k-fold cross-validation

Occam's razor in simple words states that one should not try to complicate things that can be solved in a much simpler manner. In other words, the simplest solutions are the most generalizable solutions. In general, whenever your model does not obey Occam's razor, it is probably overfitting.

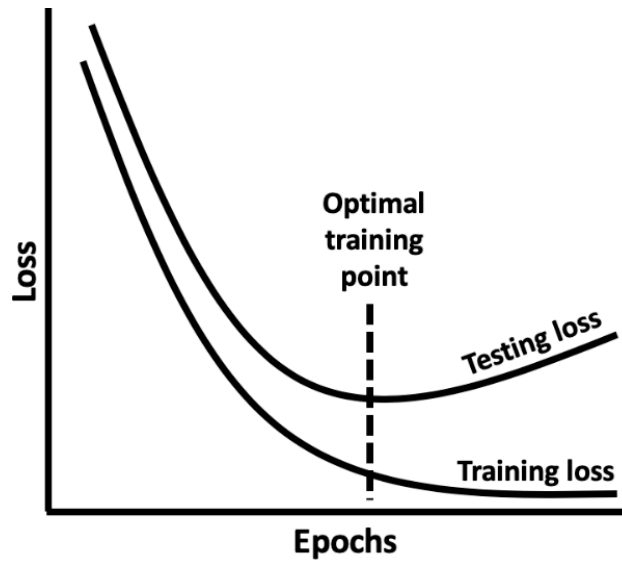


Figure 1: General Understanding for Over-fitting

Regarding stratified k-fold cross-validation, the important technique is that you will have k-different models which have different parameters, and these k-different model will be used for creating a inference which will be created by merging or ensembling from k-different inference results.

1.1.1 k-fold cross-validation

```
from sklearn.model_selection import KFold
from tabulate import tabulate
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/wine-quality.csv"
df = pd.read_csv(url, sep=";")

kf = KFold(n_splits=5, shuffle=True, random_state=42)

df['kfold'] = -1

# Split df into 5 folds
for fold, (train_index, test_index) in enumerate(kf.split(df)):
    df.loc[test_index, 'kfold'] = fold

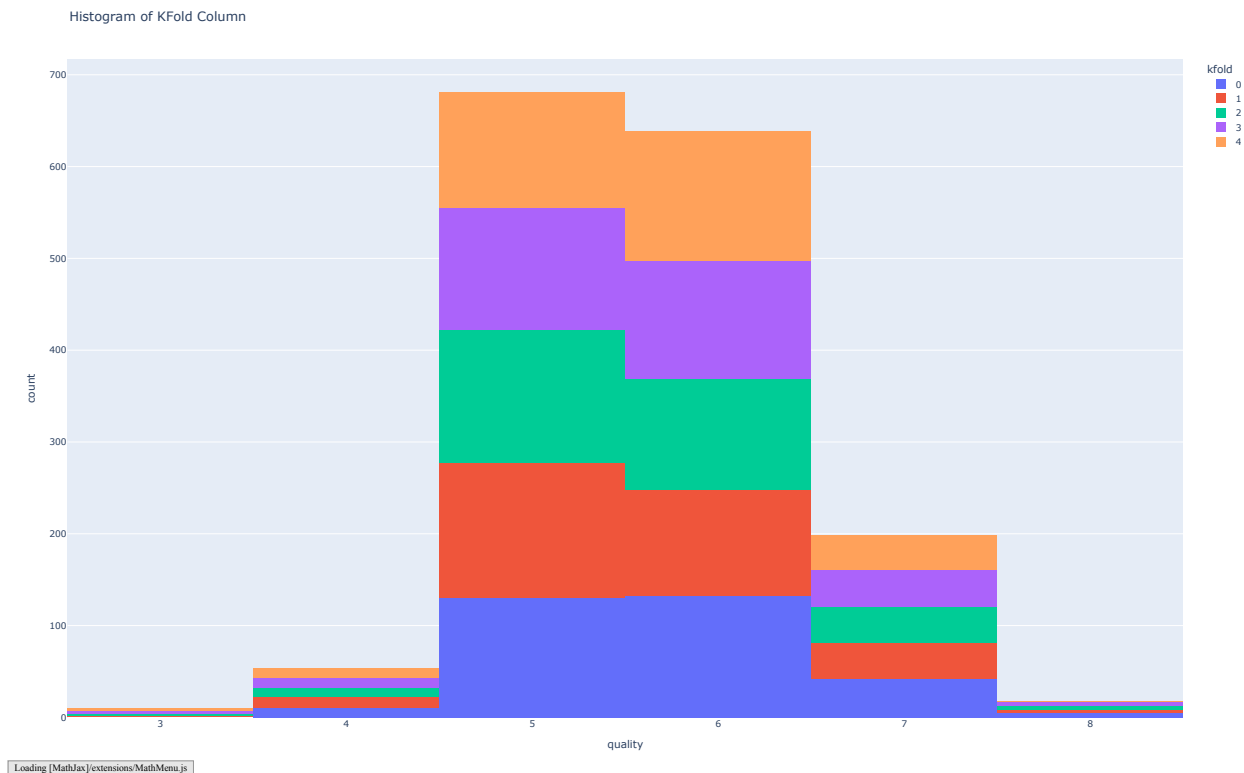
df.to_csv('../data/wine-quality-data.csv')
print(df)
```

	fixed acidity	volatile acidity	citric acid	...	alcohol	quality	kfold
0	7.4	0.700	0.00	...	9.4	5	2
1	7.8	0.880	0.00	...	9.8	5	4
2	7.8	0.760	0.04	...	9.8	5	2
3	11.2	0.280	0.56	...	9.8	6	2
4	7.4	0.700	0.00	...	9.4	5	3
...
1594	6.2	0.600	0.08	...	10.5	5	2
1595	5.9	0.550	0.10	...	11.2	6	4
1596	6.3	0.510	0.13	...	11.0	6	3
1597	5.9	0.645	0.12	...	10.2	5	1
1598	6.0	0.310	0.47	...	11.0	6	3

[1599 rows x 13 columns]

Would be much better to show as a graph as well

```
import plotly.express as px
df = df.sort_values(by="kfold")
fig = px.histogram(df, x="quality", color="kfold", title="Histogram of KFold Column")
fig.show()
```



1.1.2 stratified k-fold cross-validation

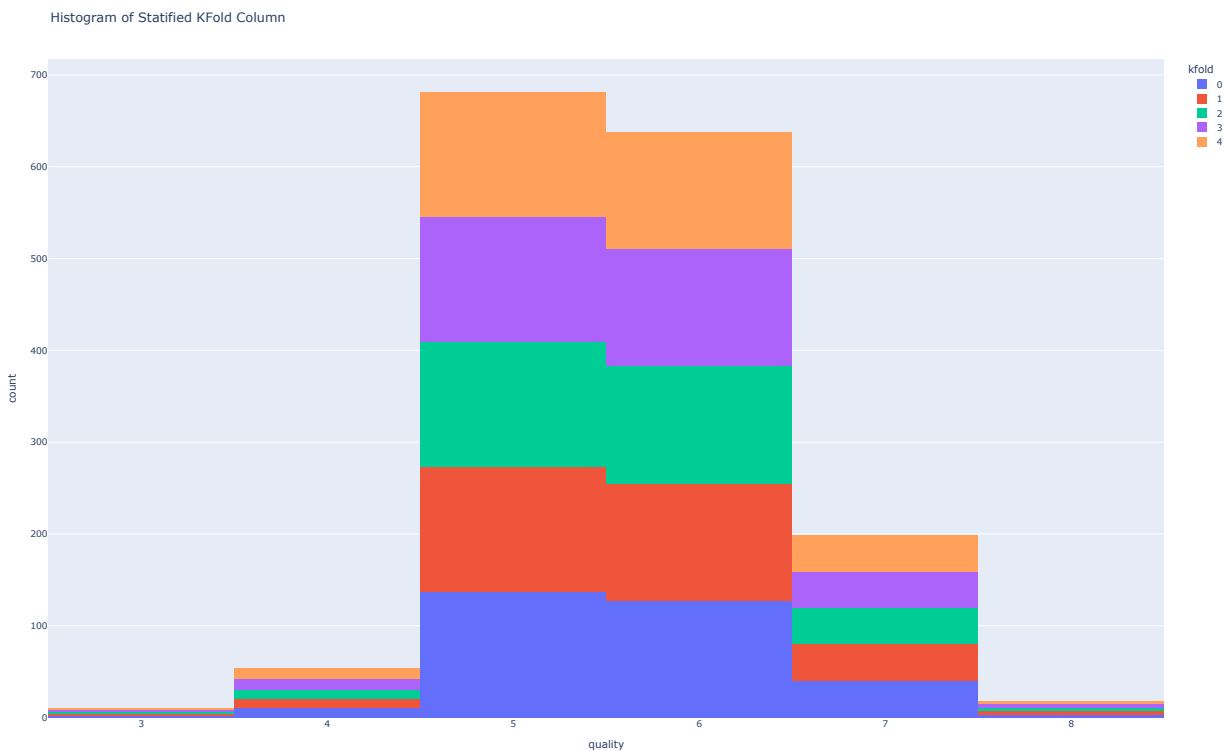
```
from sklearn.model_selection import StratifiedKFold
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-  
df = pd.read_csv(url, sep=";")
# Create a StratifiedKFold object
skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

df['kfold'] = -1

# Split df into 5 folds
for fold, (train_index, test_index) in enumerate(skf.split(df, df['quality'])):
    df.loc[test_index, 'kfold'] = fold
```

```
import plotly.express as px

df = df.sort_values(by="kfold")
fig = px.histogram(df, x="quality", color="kfold", title="Histogram of Statified  
fig.show()
```



1.1.3 hold-out based validation

In some cases, stratified k-fold cross-validation is quite demanding for computing. For this kind of case, just one fold is used for validation set. And it is recommended to split into higher number of k-fold if the number of samples is high such as 1 million.

1.2 Regression

- Mostly, simple k-fold cross-validation works for any regression problem
- If you see that the distribution of targets is not consistent, you can use stratified k-fold
- If you have a lot of samples ($> 10k$, $> 100k$), then you don't need to care about the number of bins.